

Predicting Apartment Prices

Team 6

Paula Demacker, Peter Mankiewich, (Vince)Chenzhi Pan, (Sylar)Jiajian Guo

Short Recap

Problem: How to decide on a reasonable price for an apartment?

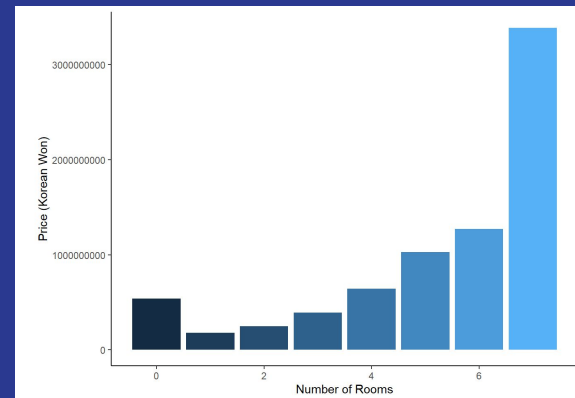
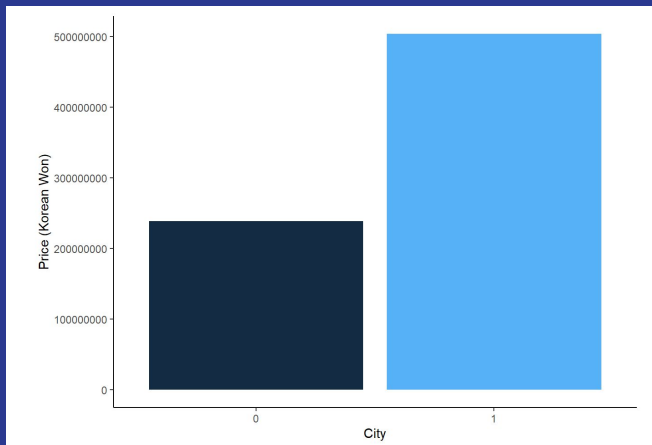
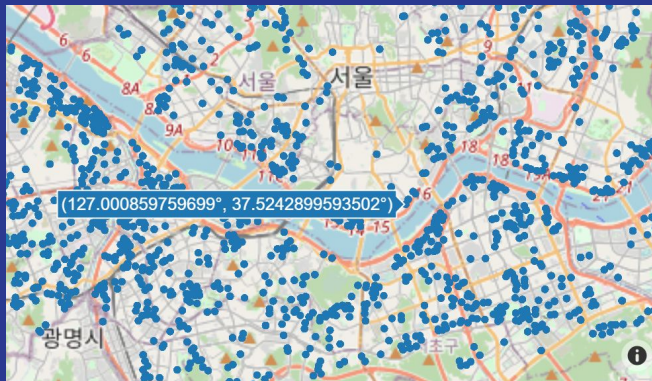
- CSV-File downloaded via Kaggle
- 1,6 Million rows
- 25 columns
- numeric, character and logical variables
- Target variable: Sales Price

Real World Application:

- Apartment buyers can compare with this model's predicted prices to estimate if their offered prices by sellers were overvalued or undervalued
- Real estate developers can predict future sales and calculate ROI (Return on Investment) more accurately before the apartments are open to market yet


EDA and Cleaning

- Developed some charts to understand the distribution of variables and their relationships
- Verified location information
- Created dummy variables for later machine learning steps



Lasso regression

- Best Lambda = 16,491
- Test RMSE = 188,678,874
- No zero coefficients
- All variables are good predictors



Main Results

- Variables importances;
Based on Tree, Boosting , consistent with each other
- OLS, Lasso, Ridge, Forward, Backward, Trees
-> better interpretation
Bagging, Boosting, RandomForest -> relatively more accurate
- Boosting

Method name	Training RMSE	Test RMSE
OLS	1.88332e+8	1.88545e+8
Forward	1.88324e+8	1.88550e+8
Backward	1.88322e+8	1.88550e+8
Ridge	2.01424e+8	2.01931e+8
Lasso	1.88455e+8	1.88679e+8
Tree	1.82321e+8	1.98231e+8
Bagging	3.6043e+7	9.13133e+7
Random Forest	4.3019e+7	9.33058e+7
Boosting	2.7564e+7	4.5081e+5
Boosting with FE	4.0315e+7	5.01138e+7

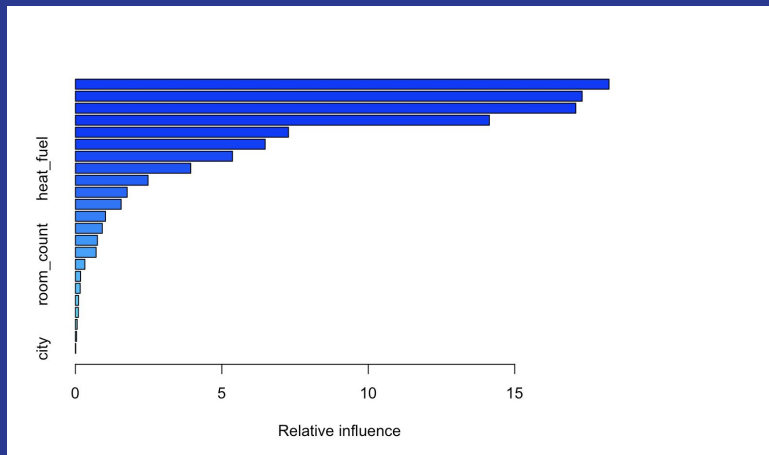
Choosing the Best Model

Boosting

- Randomly use 8000 rows
- 5000 Lowest CV error
- Use 4 as depth
- Lowest Test RMSE (8×10^5)
- Significantly better than other method
- RMSE varies a lot (Use different test sample test)


Relatively important predictors

- Supply_area
- Exclusive_use_area
- address_by_law



Feature engineering

- Converting latitude and longitude into a 3d space
 - $x: \cos(\text{longitude}) * \cos(\text{latitude})$
 - $y: \cos(\text{longitude}) * \sin(\text{latitude})$
 - $z: \sin(\text{latitude})$
- Supply area: area of entire apartment complex
- Exclusive use area: total floor area of building
- Living area = supply area - exclusive use area
- Area ratio = $\text{exclusive_use_area} / \text{supply_area}$



	Training RMSE	Test RMSE
Boosting	2.7564e+7	7.17116e+5
Boosting with FE	4.0315e+7	5.01138e+7

Challenges

- Language & Industry knowledge:
 - *Challenge:* Korean auto-translated English columns names without explanations
 - *Solution:* Google satellite images with coordinates and learned Korean apartment structures
- Data Cleaning:
 - *Challenge:* Miscellaneous character type qualitative variable inputs data
 - *Solution:* Hard coded into dummy variables or numerics
- Modeling
 - *Challenge:* Little RMSE improvement in linear models
 - *Solution:* Performed a wide range of different models with flexibilities
- Collaboration:
 - *Challenge:* Combing everyone's code into one final file without conflicts
 - *Solution:* Used Git linking RStudio and avoided accidental overwrites

Conclusion

- What We Learned:

- Needless to mess with millions rows of data when similar results can be generated by a 1/10 size
- Thought the seemingly intuitional dataset should perform well on simple models but totally the opposite
- Code collaborations suck without cloud environments
- Industry Applications:
 - Help predict prices based off of the characteristics of apartments and their immediate environment
 - Speed up real estate agents price prediction labelling work by automation

- What to Put into Practice in Future:

- Spend more time on models which perform better
- Use cloud environments for team projects
- Collect unstructured data to improve prediction accuracy if possible: satellite images for computer vision



Q&A

OLS Linear Regression

```
call:
lm(formula = transaction_real_price ~ ., data = dd_train)

Residuals:
    Min       1Q   Median       3Q      Max
-1.118e+09 -9.743e+07 -1.692e+07  7.074e+07  6.097e+09

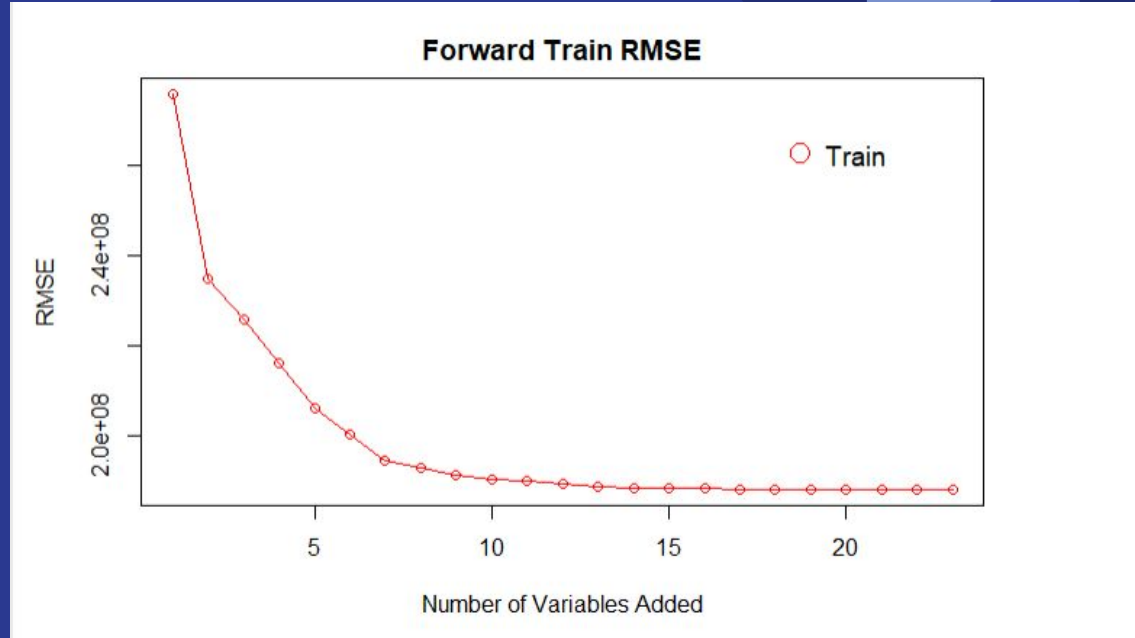
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -5.436e+10  9.421e+08  -57.707 < 2e-16 ***
city         5.765e+09  5.642e+07  102.180 < 2e-16 ***
transaction_year_month 1.668e+05  1.428e+03  116.823 < 2e-16 ***
transaction_date      6.465e+05  6.468e+05   1.000 0.317529
year_of_completion  -2.637e+06  9.746e+04  -27.053 < 2e-16 ***
exclusive_use_area    1.553e+06  1.346e+05   11.538 < 2e-16 ***
floor           1.869e+06  8.684e+04   21.525 < 2e-16 ***
latitude        -1.026e+09  1.106e+07  -92.799 < 2e-16 ***
longitude       4.498e+08  7.047e+06   63.830 < 2e-16 ***
address_by_law     1.417e+00  3.967e-02   35.713 < 2e-16 ***
total_parking_capacity_in_site 2.850e+04  1.047e+03   27.226 < 2e-16 ***
total_household_count_in_sites -8.174e+04  1.679e+03  -48.684 < 2e-16 ***
apartment_building_count_in_sites 6.668e+06  8.495e+04   78.494 < 2e-16 ***
tallest_building_in_sites 2.406e+06  1.213e+05   19.832 < 2e-16 ***
lowest_building_in_sites 3.271e+06  1.199e+05   27.291 < 2e-16 ***
heat_fuel         5.167e+07  3.737e+06   13.828 < 2e-16 ***
supply_area       3.326e+06  1.148e+05   28.977 < 2e-16 ***
total_household_count_of_area_type -1.459e+04  1.929e+03  -7.565 3.90e-14 ***
room_count       -6.117e+06  1.280e+06  -4.778 1.77e-06 ***
bathroom_count   -6.561e+06  1.725e+06  -3.803 0.000143 ***
heat_type_central  7.847e+06  2.117e+06   3.707 0.000210 ***
heat_type_district 6.760e+07  3.779e+06   17.890 < 2e-16 ***
front_door_structure_corridor -1.441e+07  1.760e+06  -8.184 2.77e-16 ***
front_door_structure_mixed -2.789e+06  4.178e+06  -0.668 0.504395
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 188300000 on 127976 degrees of freedom
Multiple R-squared:  0.6566,    Adjusted R-squared:  0.6565
F-statistic: 1.064e+04 on 23 and 127976 DF,  p-value: < 2.2e-16
```

Simple Model


- Train RMSE: 188,332,311
- Test RMSE: 188,544,540

Forward Selection



- Only intercept: train RMSE is 275,820,794
- After adding 21 predictors one by one, Train MSE kept reducing to 188,321,979.
- Adding transaction_date & front_door_structure_mixed increases RMSE
- Test RMSE is 188,550,321

Backward Selection



- Only intercept: train RMSE is 275,820,794
- After adding 21 predictors one by one, Train RMSE kept reducing to 188,321,979.
- Adding transaction_date & front_door_structure_mixed increases RMSE
- Test RMSE is 188,550,321

Ridge Regression

Train RMSE: 201425915 Test RMSE: 201936147 Best Lambda: 16491016

- Train RMSE: 201425915
- Test RMSE: 201936147
- Best Lambda: 16491016

Bagging

```
[35]: set.seed(217)
      test.mse = c()

      for (i in seq(1,5)) {
        train = sample(1:nrow(Price),(nrow(Price)/100)*i)
        tree.testy = Price[-train,transaction_real_price]
        tree.test = Price[-train]
        bag.price = randomForest(transaction_real_price ~ ., data = Price, subset = train, mtry = 24, importance = TRUE)
        yhat.bag = predict(bag.price, newdata = tree.test)
        test.mse = c(test.mse,mean((tree.test - yhat.bag)^2))
      }

      test.mse
```

**MORE THAN 12 HOURS to train in
GOOGLE AI NOTEBOOK with
EXPENSIVE COMPUTER OPTION**

6538671859298821 · 4728267858107837 · 4151840578511839 · 3401919969449496 · 3153720272922860

Decision Tree

